

# COMPARING REGRESSION MODELS

Hello everyone. This is a notebook comparing various regression models such as Ridge, Knn, Bayesian Regression, Decision Tree and SVM. It is extremely beneficial for beginners to take a close look at the notebook so as to get an insight as to how different algorithms work and also which algorithms can perform better in some cases depending upon cases

```
[1]: # This Python 3 environment comes with many helpful analytics libraries
      ↳ installed
      # It is defined by the kaggle/python docker image: https://github.com/kaggle/
      ↳ docker-python
      # For example, here's several helpful packages to load in

import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)

# Input data files are available in the "../input/" directory.
# For example, running this (by clicking run or pressing Shift+Enter) will list
↳ the files in the input directory

from subprocess import check_output
print(check_output(["ls", "../input"]).decode("utf8"))

# Any results you write to the current directory are saved as output.
```

movie\_metadata.csv

```
[2]: # Importing packages

import os
import pandas as pd
from pandas import DataFrame, Series
from sklearn import tree
import matplotlib
import numpy as np
import matplotlib.pyplot as plt
from sklearn import svm
from sklearn.preprocessing import StandardScaler
import statsmodels.formula.api as smf
```

```
import statsmodels.api as sm
from mpl_toolkits.mplot3d import Axes3D
import seaborn as sns
from sklearn import neighbors
from sklearn import linear_model
%matplotlib inline
```

```
[3]: f = pd.read_csv("../input/movie_metadata.csv")
```

```
[4]: data=DataFrame(f)
data.head()[:2]
```

```
[4]:   color  director_name  num_critic_for_reviews  duration \
0  Color   James Cameron                723.0    178.0
1  Color   Gore Verbinski                302.0    169.0

   director_facebook_likes  actor_3_facebook_likes   actor_2_name \
0                      0.0                 855.0  Joel David Moore
1                 563.0                 1000.0   Orlando Bloom

   actor_1_facebook_likes   gross   genres \
0                 1000.0  760505847.0  Action|Adventure|Fantasy|Sci-Fi
1                40000.0  309404152.0      Action|Adventure|Fantasy

   ...   num_user_for_reviews  language  country  content_rating \
0   ...                 3054.0   English    USA          PG-13
1   ...                 1238.0   English    USA          PG-13

   budget  title_year  actor_2_facebook_likes  imdb_score  aspect_ratio \
0  237000000.0    2009.0                 936.0         7.9         1.78
1  300000000.0    2007.0                 5000.0         7.1         2.35

   movie_facebook_likes
0                 33000
1                      0
```

[2 rows x 28 columns]

*Getting non-object elements*

```
[5]: X_data=data.dtypes[data.dtypes!='object'].index
X_train=data[X_data]
X_train.head()[:2]
```

```
[5]:   num_critic_for_reviews  duration  director_facebook_likes \
0                723.0    178.0                 0.0
1                302.0    169.0                563.0
```

	actor_3_facebook_likes	actor_1_facebook_likes	gross	\
0	855.0	1000.0	760505847.0	
1	1000.0	40000.0	309404152.0	

	num_voted_users	cast_total_facebook_likes	facenumber_in_poster	\
0	886204	4834	0.0	
1	471220	48350	0.0	

	num_user_for_reviews	budget	title_year	actor_2_facebook_likes	\
0	3054.0	237000000.0	2009.0	936.0	
1	1238.0	300000000.0	2007.0	5000.0	

	imdb_score	aspect_ratio	movie_facebook_likes
0	7.9	1.78	33000
1	7.1	2.35	0

```
[6]: X_train.describe()
```

```
[6]:
```

	num_critic_for_reviews	duration	director_facebook_likes	\
count	4993.000000	5028.000000	4939.000000	
mean	140.194272	107.201074	686.509212	
std	121.601675	25.197441	2813.328607	
min	1.000000	7.000000	0.000000	
25%	50.000000	93.000000	7.000000	
50%	110.000000	103.000000	49.000000	
75%	195.000000	118.000000	194.500000	
max	813.000000	511.000000	23000.000000	

	actor_3_facebook_likes	actor_1_facebook_likes	gross	\
count	5020.000000	5036.000000	4.159000e+03	
mean	645.009761	6560.047061	4.846841e+07	
std	1665.041728	15020.759120	6.845299e+07	
min	0.000000	0.000000	1.620000e+02	
25%	133.000000	614.000000	5.340988e+06	
50%	371.500000	988.000000	2.551750e+07	
75%	636.000000	11000.000000	6.230944e+07	
max	23000.000000	640000.000000	7.605058e+08	

	num_voted_users	cast_total_facebook_likes	facenumber_in_poster	\
count	5.043000e+03	5043.000000	5030.000000	
mean	8.366816e+04	9699.063851	1.371173	
std	1.384853e+05	18163.799124	2.013576	
min	5.000000e+00	0.000000	0.000000	
25%	8.593500e+03	1411.000000	0.000000	
50%	3.435900e+04	3090.000000	1.000000	
75%	9.630900e+04	13756.500000	2.000000	

max	1.689764e+06	656730.000000	43.000000
-----	--------------	---------------	-----------

	num_user_for_reviews	budget	title_year \
count	5022.000000	4.551000e+03	4935.000000
mean	272.770808	3.975262e+07	2002.470517
std	377.982886	2.061149e+08	12.474599
min	1.000000	2.180000e+02	1916.000000
25%	65.000000	6.000000e+06	1999.000000
50%	156.000000	2.000000e+07	2005.000000
75%	326.000000	4.500000e+07	2011.000000
max	5060.000000	1.221550e+10	2016.000000

	actor_2_facebook_likes	imdb_score	aspect_ratio	movie_facebook_likes
count	5030.000000	5043.000000	4714.000000	5043.000000
mean	1651.754473	6.442138	2.220403	7525.964505
std	4042.438863	1.125116	1.385113	19320.445110
min	0.000000	1.600000	1.180000	0.000000
25%	281.000000	5.800000	1.850000	0.000000
50%	595.000000	6.600000	2.350000	166.000000
75%	918.000000	7.200000	2.350000	3000.000000
max	137000.000000	9.500000	16.000000	349000.000000

```
[7]: # Finding all the columns with NULL values
```

```
np.sum(X_train.isnull())
```

```
[7]: num_critic_for_reviews    50
duration                    15
director_facebook_likes    104
actor_3_facebook_likes     23
actor_1_facebook_likes      7
gross                      884
num_voted_users             0
cast_total_facebook_likes   0
facenumber_in_poster       13
num_user_for_reviews        21
budget                     492
title_year                 108
actor_2_facebook_likes     13
imdb_score                  0
aspect_ratio               329
movie_facebook_likes        0
dtype: int64
```

```
[8]: # Filling all Null values
```

```
X_train=X_train.fillna(0)
columns=X_train.columns.tolist()
```

```
y=X_train['imdb_score']
X_train.drop(['imdb_score'],axis=1,inplace=True)
X_train.head()[:2]
```

```
[8]:   num_critic_for_reviews  duration  director_facebook_likes  \
0                723.0      178.0                0.0
1                302.0      169.0               563.0

      actor_3_facebook_likes  actor_1_facebook_likes      gross  \
0                855.0                1000.0  760505847.0
1               1000.0               40000.0  309404152.0

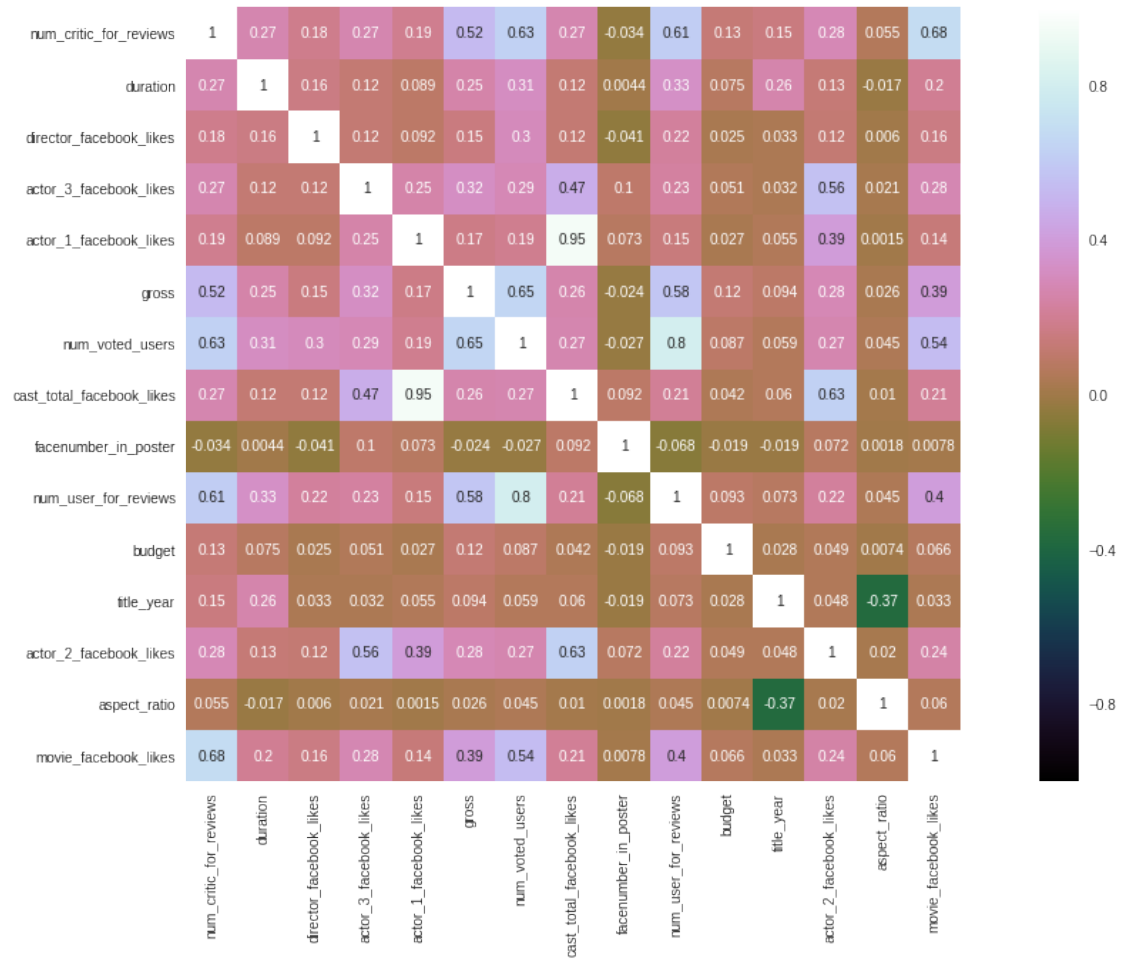
      num_voted_users  cast_total_facebook_likes  facenumber_in_poster  \
0             886204                4834                0.0
1             471220                48350                0.0

      num_user_for_reviews      budget  title_year  actor_2_facebook_likes  \
0                3054.0  237000000.0      2009.0                936.0
1                1238.0  300000000.0      2007.0               5000.0

      aspect_ratio  movie_facebook_likes
0             1.78                33000
1             2.35                 0
```

```
[9]: # GETTING Correllation matrix
corr_mat=X_train.corr(method='pearson')
plt.figure(figsize=(20,10))
sns.heatmap(corr_mat,vmax=1,square=True,annot=True,cmap='cubehelix')
```

```
[9]: <matplotlib.axes._subplots.AxesSubplot at 0x7f3c0fa12cc0>
```

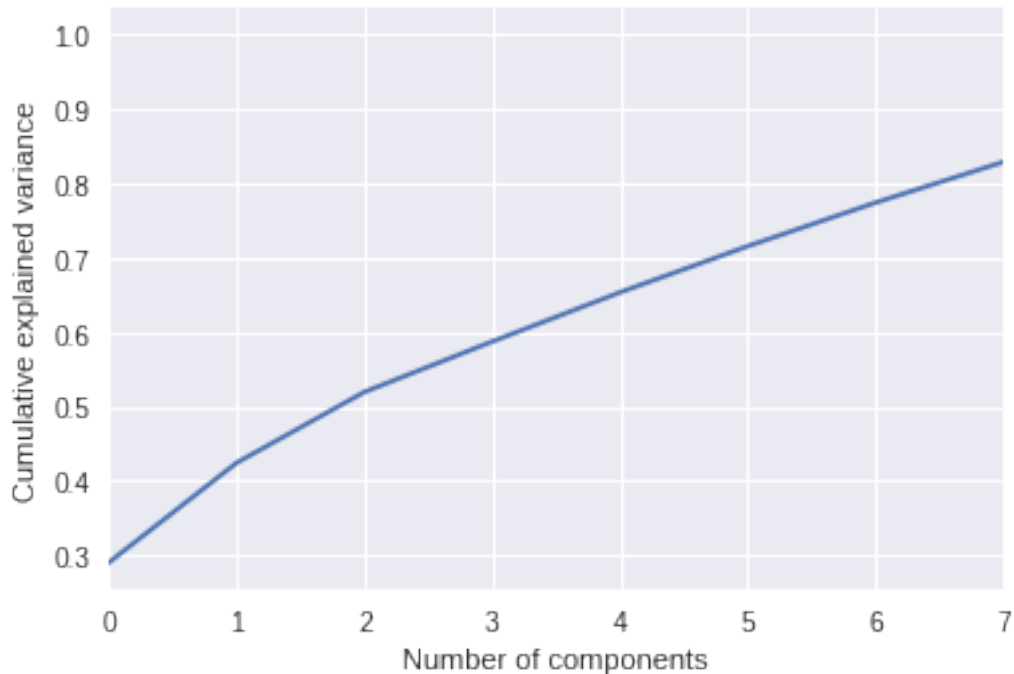


```
[10]: X_Train=X_train.values
X_Train=np.asarray(X_Train)

# Finding normalised array of X_Train
X_std=StandardScaler().fit_transform(X_Train)
```

```
[11]: from sklearn.decomposition import PCA
pca = PCA().fit(X_std)
plt.plot(np.cumsum(pca.explained_variance_ratio_))
plt.xlim(0,7,1)
plt.xlabel('Number of components')
plt.ylabel('Cumulative explained variance')
```

```
[11]: <matplotlib.text.Text at 0x7f3c05789278>
```

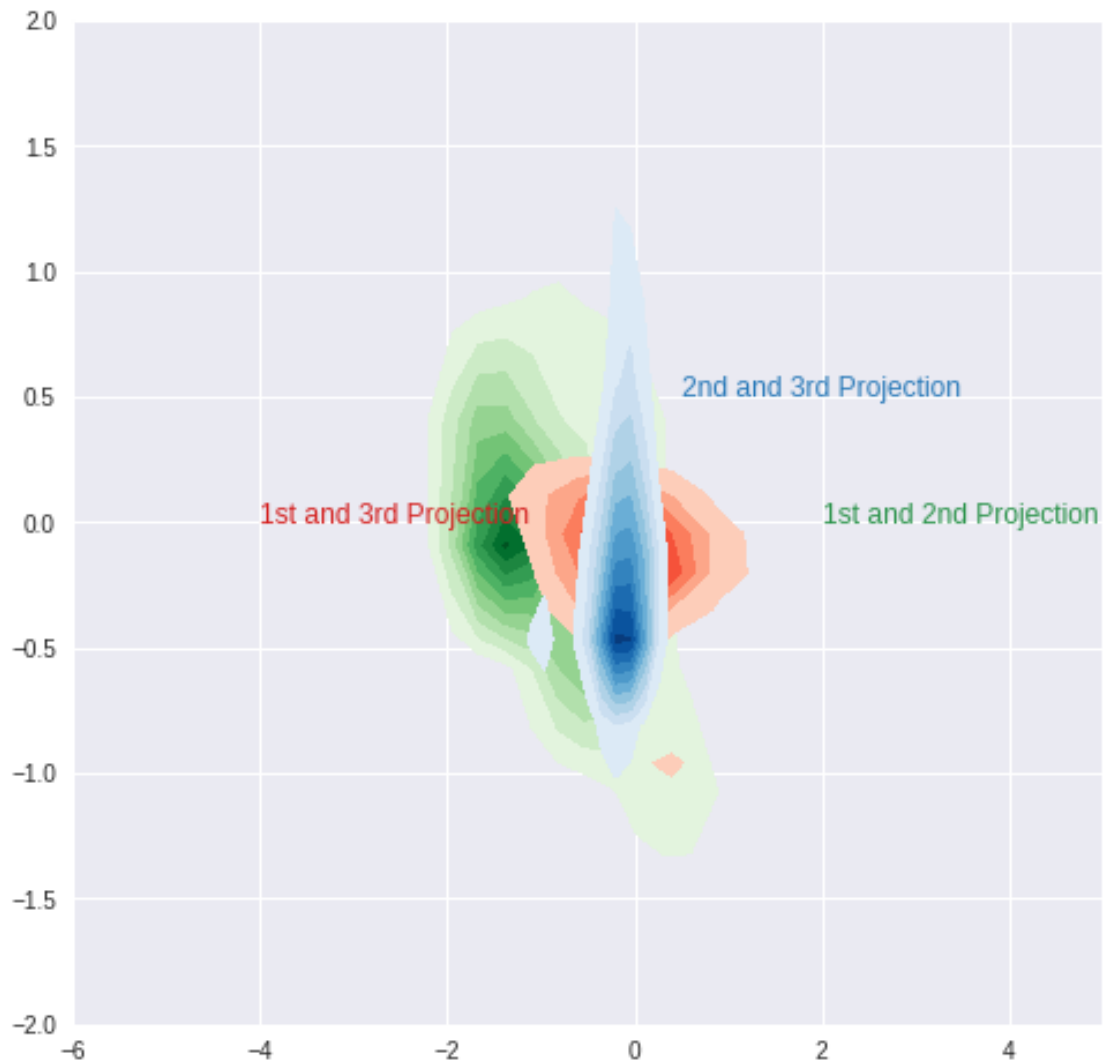


Since 5 components can explain more than 70% of the variance, we choose the number of the components to be 5

```
[12]: from sklearn.decomposition import PCA
sklearn_pca=PCA(n_components=5)
X_Train=sklearn_pca.fit_transform(X_std)

sns.set(style='darkgrid')
f, ax = plt.subplots(figsize=(8, 8))
# ax.set_aspect('equal')
ax = sns.kdeplot(X_Train[:,0], X_Train[:,1], cmap="Greens",
                shade=True, shade_lowest=False)
ax = sns.kdeplot(X_Train[:,1], X_Train[:,2], cmap="Reds",
                shade=True, shade_lowest=False)
ax = sns.kdeplot(X_Train[:,2], X_Train[:,3], cmap="Blues",
                shade=True, shade_lowest=False)
red = sns.color_palette("Reds")[-2]
blue = sns.color_palette("Blues")[-2]
green = sns.color_palette("Greens")[-2]
ax.text(0.5, 0.5, "2nd and 3rd Projection", size=12, color=blue)
ax.text(-4, 0.0, "1st and 3rd Projection", size=12, color=red)
ax.text(2, 0, "1st and 2nd Projection", size=12, color=green)
plt.xlim(-6,5)
plt.ylim(-2,2)
```

[12]: (-2, 2)



```
[13]: number_of_samples = len(y)
      np.random.seed(0)
      random_indices = np.random.permutation(number_of_samples)
      num_training_samples = int(number_of_samples*0.75)
      x_train = X_Train[random_indices[:num_training_samples]]
      y_train=y[random_indices[:num_training_samples]]
      x_test=X_Train[random_indices[num_training_samples:]]
      y_test=y[random_indices[num_training_samples:]]
      y_Train=list(y_train)
```

## Ridge Regression



```
[14]: model=linear_model.Ridge()
model.fit(x_train,y_train)
y_predict=model.predict(x_train)

error=0
for i in range(len(y_Train)):
    error+=(abs(y_Train[i]-y_predict[i])/y_Train[i])
train_error_ridge=error/len(y_Train)*100
print("Train error = '{}'".format(train_error_ridge)+" percent in Ridge_
↪Regression")

Y_test=model.predict(x_test)
y_Predict=list(y_test)

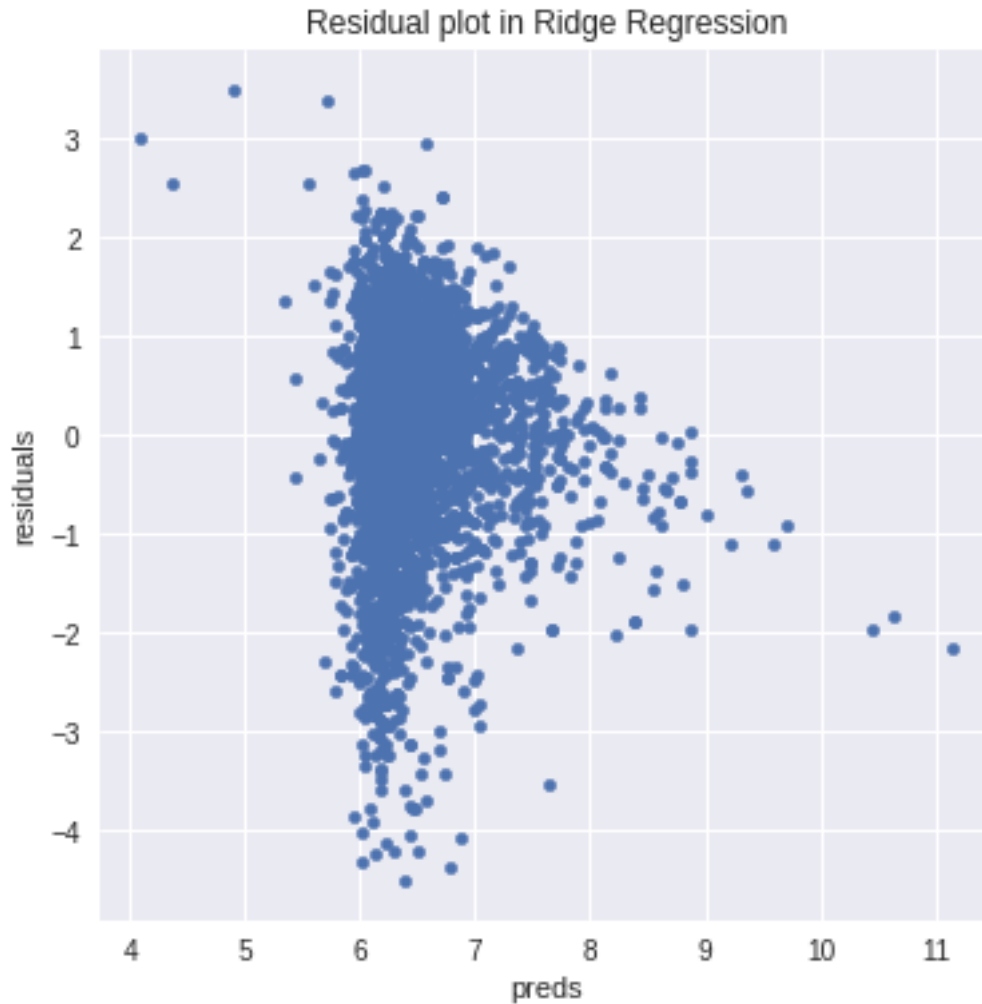
error=0
for i in range(len(y_test)):
    error+=(abs(y_Predict[i]-Y_test[i])/y_Predict[i])
test_error_ridge=error/len(Y_test)*100
print("Test error = '{}'".format(test_error_ridge)+" percent in Ridge_
↪Regression")
```

Train error = 13.914226734021002 percent in Ridge Regression  
Test error = 15.299716605526257 percent in Ridge Regression

```
[15]: matplotlib.rcParams['figure.figsize'] = (6.0, 6.0)

preds = pd.DataFrame({"preds":model.predict(x_train), "true":y_train})
preds["residuals"] = preds["true"] - preds["preds"]
preds.plot(x = "preds", y = "residuals",kind = "scatter")
plt.title("Residual plot in Ridge Regression")
```

```
[15]: <matplotlib.text.Text at 0x7f3c040fd0f0>
```



## Knn Algorithm

```
[16]: n_neighbors=5
knn=neighbors.KNeighborsRegressor(n_neighbors,weights='uniform')
knn.fit(x_train,y_train)
y1_knn=knn.predict(x_train)
y1_knn=list(y1_knn)

error=0
for i in range(len(y_train)):
    error+=(abs(y1_knn[i]-y_train[i])/y_train[i])
train_error_knn=error/len(y_train)*100
print("Train error = "+'{}'.format(train_error_knn)+" percent"+" in Knn_
    ↪algorithm")

y2_knn=knn.predict(x_test)
```

```

y2_knn=list(y2_knn)
error=0
for i in range(len(y_test)):
    error+=(abs(y2_knn[i]-Y_test[i])/Y_test[i])
test_error_knn=error/len(Y_test)*100
print("Test error = '{}''.format(test_error_knn)+" percent"+" in knn algorithm")

```

Train error = 10.812937212714084 percent in Knn algorithm

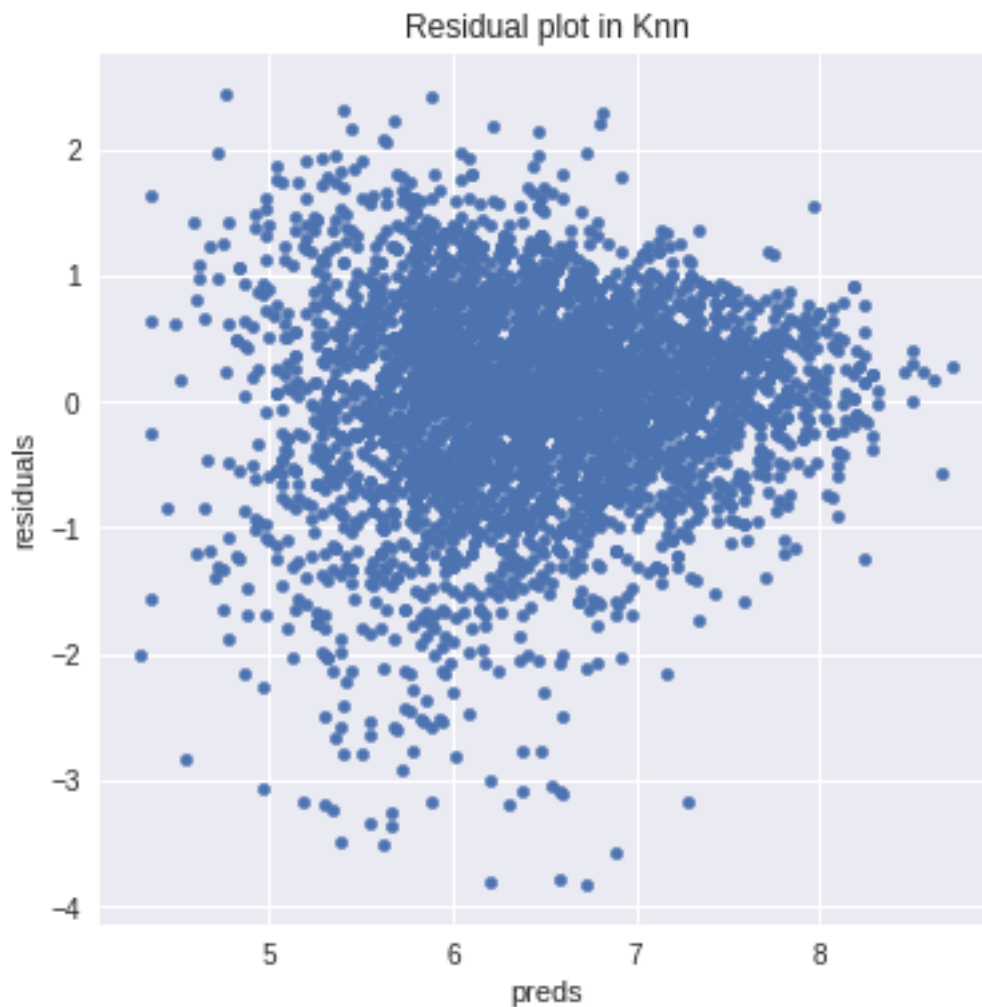
Test error = 6.878221673331934 percent in knn algorithm

```

[17]: matplotlib.rcParams['figure.figsize'] = (6.0, 6.0)
preds = pd.DataFrame({"preds":knn.predict(x_train), "true":y_train})
preds["residuals"] = preds["true"] - preds["preds"]
preds.plot(x = "preds", y = "residuals",kind = "scatter")
plt.title("Residual plot in Knn")

```

[17]: <matplotlib.text.Text at 0x7f3bfc306160>



## Bayesian Regression

```
[18]: reg = linear_model.BayesianRidge()
reg.fit(x_train,y_train)
y1_reg=reg.predict(x_train)
y1_reg=list(y1_reg)
y2_reg=reg.predict(x_test)
y2_reg=list(y2_reg)

error=0
for i in range(len(y_train)):
    error+=(abs(y1_reg[i]-y_Train[i])/y_Train[i])
train_error_bay=error/len(y_Train)*100
print("Train error = "+'{}'.format(train_error_bay)+" percent"+" in Bayesian_
↳Regression")

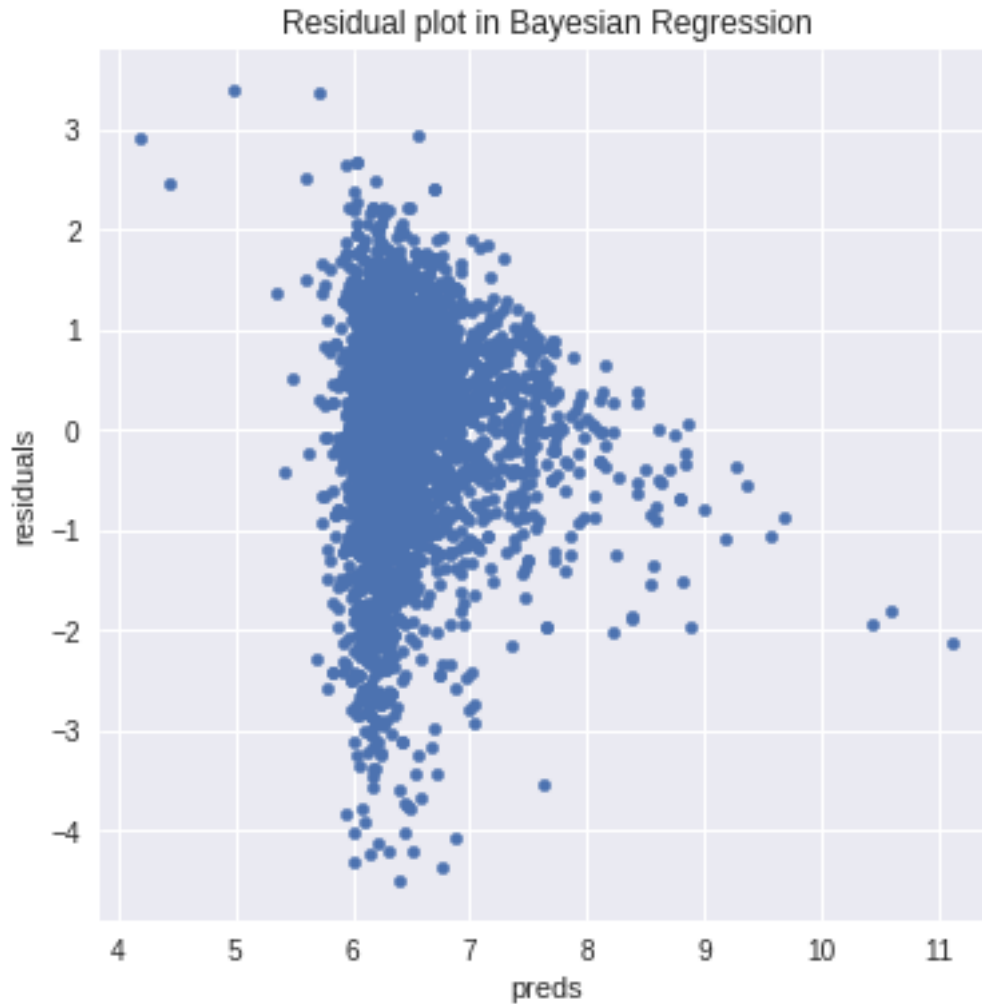
error=0
for i in range(len(y_test)):
    error+=(abs(y2_reg[i]-Y_test[i])/Y_test[i])
test_error_bay=(error/len(Y_test))*100
print("Test error = "+'{}'.format(test_error_bay)+" percent"+" in Bayesian_
↳Regression")
```

Train error = 13.91749661366315 percent in Bayesian Regression

Test error = 0.025287435537397897 percent in Bayesian Regression

```
[19]: matplotlib.rcParams['figure.figsize'] = (6.0, 6.0)
preds = pd.DataFrame({"preds":reg.predict(x_train), "true":y_train})
preds["residuals"] = preds["true"] - preds["preds"]
preds.plot(x = "preds", y = "residuals",kind = "scatter")
plt.title("Residual plot in Bayesian Regression")
```

```
[19]: <matplotlib.text.Text at 0x7f3bfc2f80b8>
```



### Decision Tree Regressor

```
[20]: dec = tree.DecisionTreeRegressor(max_depth=1)
dec.fit(x_train,y_train)
y1_dec=dec.predict(x_train)
y1_dec=list(y1_dec)
y2_dec=dec.predict(x_test)
y2_dec=list(y2_dec)

error=0
for i in range(len(y_train)):
    error+=(abs(y1_dec[i]-y_Train[i])/y_Train[i])
train_error_tree=error/len(y_Train)*100
print("Train error = "+'{}'.format(train_error_tree)+" percent"+" in Decision_
Tree Regressor")
```

```

error=0
for i in range(len(y_test)):
    error+=(abs(y1_dec[i]-Y_test[i])/Y_test[i])
test_error_tree=error/len(Y_test)*100
print("Test error = '{}'".format(test_error_tree)+" percent in Decision Tree_
↪Regressor")

```

Train error = 14.590941891509965 percent in Decision Tree Regressor


Test error = 5.816650087351861 percent in Decision Tree Regressor

```

[21]: matplotlib.rcParams['figure.figsize'] = (6.0, 6.0)
preds = pd.DataFrame({"preds":dec.predict(x_train), "true":y_train})
preds["residuals"] = preds["true"] - preds["preds"]
preds.plot(x = "preds", y = "residuals",kind = "scatter")
plt.title("Residual plot in Decision Tree")

```

[21]: <matplotlib.text.Text at 0x7f3bfc22f7b8>



## SVM

```
[22]: svm_reg=svm.SVR()
      svm_reg.fit(x_train,y_train)
      y1_svm=svm_reg.predict(x_train)
      y1_svm=list(y1_svm)
      y2_svm=svm_reg.predict(x_test)
      y2_svm=list(y2_svm)

      error=0
      for i in range(len(y_train)):
          error+=(abs(y1_svm[i]-y_Train[i])/y_Train[i])
      train_error_svm=error/len(y_Train)*100
      print("Train error = "+'{}'.format(train_error_svm)+" percent"+" in SVM_
      ↪Regressor")

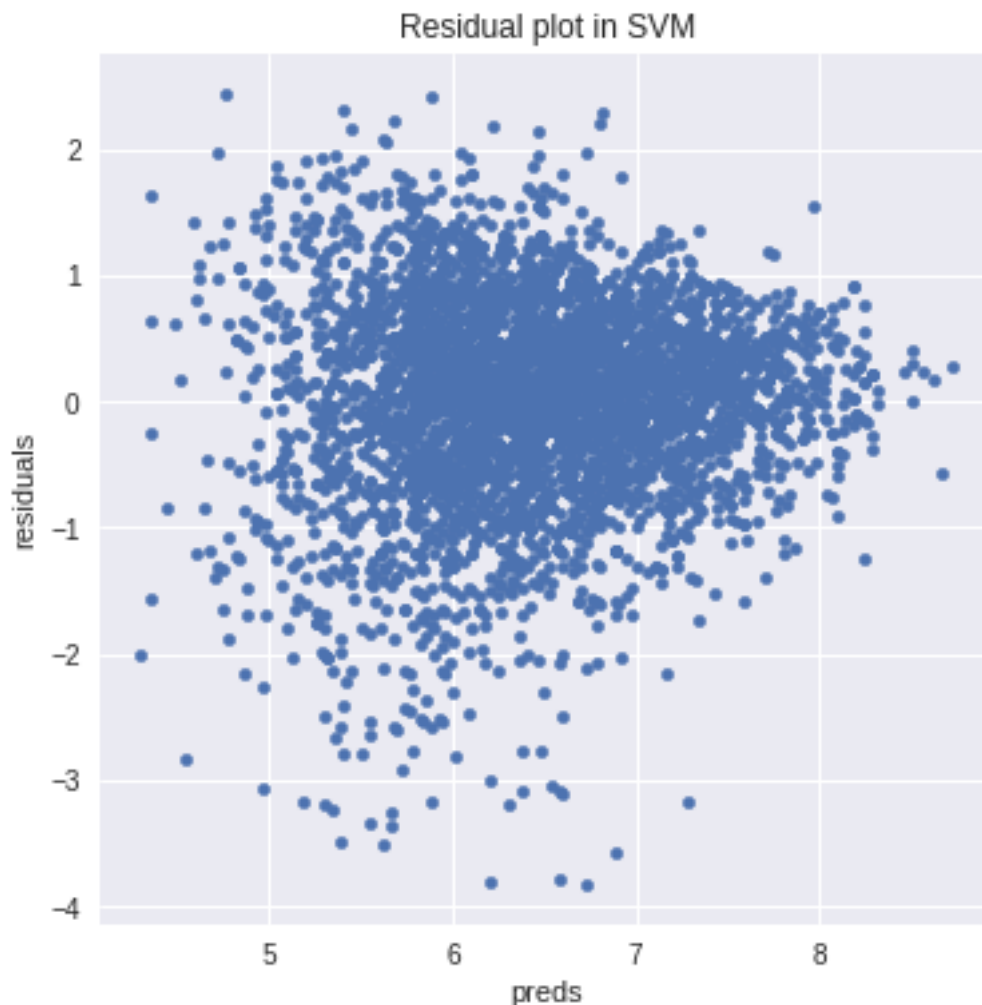
      error=0
      for i in range(len(y_test)):
          error+=(abs(y2_svm[i]-Y_test[i])/Y_test[i])
      test_error_svm=error/len(Y_test)*100
      print("Test error = '"+'{}'.format(test_error_svm)+" percent in SVM Regressor")
```

Train error = 12.036337747988636 percent in SVM Regressor

Test error = 5.403852057483367 percent in SVM Regressor

```
[23]: matplotlib.rcParams['figure.figsize'] = (6.0, 6.0)
      preds = pd.DataFrame({"preds":knn.predict(x_train), "true":y_train})
      preds["residuals"] = preds["true"] - preds["preds"]
      preds.plot(x = "preds", y = "residuals",kind = "scatter")
      plt.title("Residual plot in SVM")
```

```
[23]: <matplotlib.text.Text at 0x7f3bfc172198>
```



```
[24]: train_error=[train_error_ridge,train_error_knn,train_error_bay,train_error_tree,train_error_svm]
      test_error=[test_error_ridge,test_error_knn,test_error_bay,test_error_tree,test_error_svm]

      col={'Train Error':train_error,'Test Error':test_error}
      models=['Ridge Regression','Knn','Bayesian Regression','Decision Tree','SVM']
      df=DataFrame(data=col,index=models)
      df
```

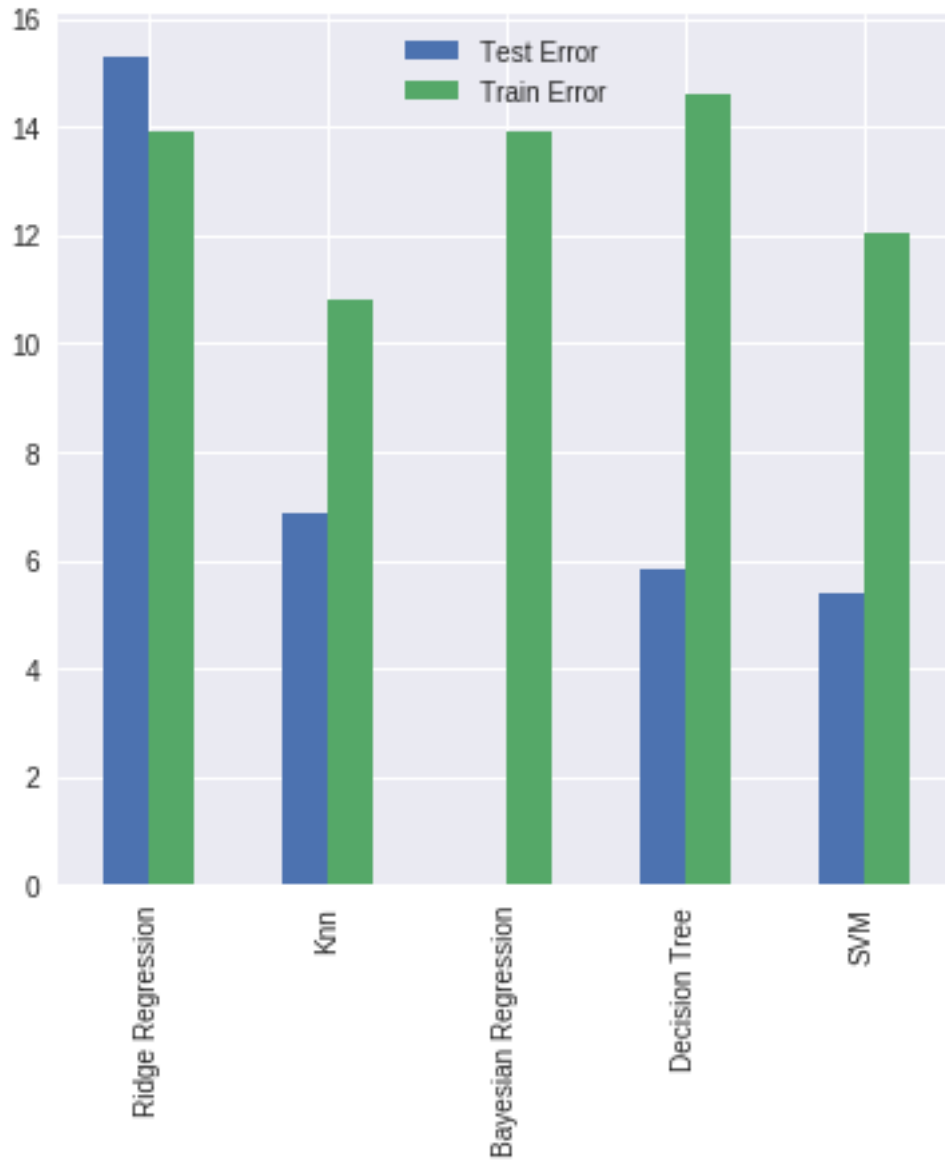
```
[24]:
```

	Test Error	Train Error
Ridge Regression	15.299717	13.914227
Knn	6.878222	10.812937
Bayesian Regression	0.025287	13.917497
Decision Tree	5.816650	14.590942
SVM	5.403852	12.036338



```
[25]: df.plot(kind='bar')
```

```
[25]: <matplotlib.axes._subplots.AxesSubplot at 0x7f3bfc0af160>
```



Seems that KNN turned out to be the winner. Its because of the fact that there are very large number of data points and also features are highly continuous. Moreover the dimensionality of the processed data is not too high.